

Natural Language Processing Homework 7

Katie Chang + Willis Wang

April 30, 2016

README

1 Q2

1.1 first

First accepts a language that starts with at least one 0 and ends with at least three 1s. There can be any combination of 0s and 1s in between the leading 0 and the trailing 111. Since the ? denotes optionality, the fourth 1 is not necessary for the string to be accepted in this language.

ii. This language is a subset of the language Bit^* .

The 0 and 1 are quoted in the .grm file because these two FSMs are defining specific strings that this language will use. Essentially, the quotes indicate the initializing of the letters of the alphabet that this language will use.

iii. There are 5 states and 9 arcs in this FST.

1.2 b. second

i. `export Second = Optimize[Zero Bit* One One One];`

ii. If First and Second are equivalent, then Disagreements should accept nothing.

Running `fstinfo` on `Disagreements.fst`, we can see that this FST has 0 states and 0 arcs, and in fact has a 0 count of everything. The initial state value is -1. Therefore, it should be unable to accept any sort of string.

1.3 c

i. First.fst now has 20 states and 25 arcs. Second.fst now has 13 states and 17 arcs. Disagreements.fst has 68 states and 88 arcs (!!!).

ii. The two separate sub-FSAs is because the Disagreement FST is taking the union of two FSTs, First - Second and Second - First.

iii. The results are the same. The only difference there is is the size of the FST and topology.

1.4 d

After creating Disagreements.fst, the number of states and arcs is the same as before. This shows that it doesn't matter whether or not we optimize the sub-FSTs, but there would be a lot of redundant work done. We should optimize if we want to avoid excessive overhead costs.

2 Q3

See binary.grm.

3 Q4

3.1 a

Input language: "a"("b"* — "c"+ — "")"a"

3.2 b

"abca" has no outputs.

"aba" has 1 output = "axa"

"aa" has 2 outputs = "aa" and "africa"

"aca" has more than 2 outputs.

3.3 c

Takes all strings that start and end with one a, as well as one of the following criteria:

1. replaces all "b"s between the two "a"s with "x"s
2. replaces all "c"s between the "a"s with any number of "y"s
3. replaces any "" between the "a"s with either the empty string or "fric".

3.4 d

Consistent. 10 states and 16 arcs.

4 5

4.1 c. Parity1

This particular transducer does not take into account the empty string.

4.2 f. UnParity

UnParity inverts the transduction. The only acceptable strings are "0" and "1" / any number of 0s before one "1". The former refers to even binary numbers, and the output lists a bunch of random ones. The latter refers to odd binary numbers, and will list a bunch of them.

5 6

5.1 a

i. With the given NP definition, the issue is that the strings are not in quotes. The program throws an error "Undefined symbol: Art". What is accepted: The list of strings that optionally start with "Art" or "Quant", has 0 or more instances of "Adj", and 1 or more instance of "Noun".

ii. There are 13 states and 17 arcs. "Adj" is represented as three separate letters in the FST, and not as a unit.

5.2 b

```
export MakeNmod = CDRewrite["Noun":"Nmod", "", "Noun", SigmaStar,
'sim', 'obl'];
```

5.3 c

i. This composition accepts only NP strings, and of those input strings, converts them with the MakeNmod rewrite FST.

ii. ArtAdjNounNounNoun -_i ArtAdjNmodNmodNoun

AdjNounNounNounVerb -_i Rewrite failed.

The second input failed because the input string is not in the NP form of (Adj*Noun+).

iii. TransformNP has 16 states and 20 arcs. MakeNmod has 36 states and 2331 arcs. This makes sense, since TransformNP accepts a smaller set of strings.

iv. The two FSTs are similar. At the end of these FSTs, NP accepts at least one "Noun". TransformNP instead takes these "Noun"s and turns them into "Nmod", except for the last one.

5.4 d

Bracket1 divides NP chunks in more ways than Bracket2 does. It will output all possible NP chunks. Bracket2 will simply go as far as possible to find the longest NP chunk.

5.5 e

```
export BracketTransform = CDRewrite[BracketNP @ MakeNmod, "", "",  
SigmaStar, 'sim', 'obl'];
```

5.6 f

See chunker.grm.

6 7

6.1 a

Check out stress.grm!

6.2 b

Input string: ev'apor'ating Output string: ev'aporating Output string: ev'apor'ating

Input string: 'incomm'unic'ado Output string: 'incommunic'ado Output string: 'incomm'unic'ado Output string: incomm'unicado Output string: 'incommunicado Output string: 'incomm'unicado Output string: incommuni-cado Output string: incommunic'ado Output string: incomm'unic'ado

6.3 c

see stress.grm.

6.4 d

Making some assumptions, we assume that if "y" is followed by a consonant, then it is probably a vowel. If it is followed by a vowel, then it is probably a consonant.

7 8

7.1 a

Lead as a noun and read as a verb. They both end with the same ending, but don't sound the same.

7.2 b

For "academic", the mouth opens more and in a more wide sense. The tongue seems to remain low in the mouth for this "a" sound.

For "academy", the mouth opens less. The sound made is more like an "uh" sound, which means the tongue doesn't really move until the "d" sound.

7.3 c

Results language: the words printed are words that start with a strong sound and go on to end with a weaker sound. A Dacytl is when a long syllable is followed by two short syllables.

7.4 d

See dactyls.grm.

7.5 f

Domain is the left hand side of the dictionary (cmudict.txt) - the words.

The range is the collection of syllables on the right side of the dictionary - the pronunciations of each syllable.

7.6 g

WordEnding @ Invert[WordEnding] gives all words with the same ending as a given input word. We can input x into WordEnding, get output y, and then use y as input into Invert[WordEnding] to get all the outputs z such that it would have the same word ending as x.

Input alphabet: Byte, Arpa Output: Byte

7.7 h

Memory allocation failed error arises. Trying to make a web between every single word in the dictionary results in a large amount of overhead.

7.8 i

This is more efficient because instead of composing the entire two FSTs together at once, we use one FST at a time. We basically make the roadmap through the FSTs rather than try to find our way through the FSTs while in the midst of it.

orange returns "orange" as the only rhyme. adventureland as a word isn't in the dictionary, and thus says "rewrite failed". But "adventure" and "land" separately have a lot of rhymes.

8 Q9

8.1 a

i. The min weight would be 1.7, and the string could be either '001' or '011'.

- ii. (001, 1011) or (011, 1111) with a weight of 1.95.

8.2 b

see binary.grm

8.3 c

i.

(x,y) = (1100, 0000)

export WeightedMultipath = (One : Zero ;2;) * (One : Zero ;1;) * (Zero ;0.1;)+;

Input string: 1100 Output string: 0000 ;cost=2.2; Output string: 0000 ;cost=3.2; Output string: 0000 ;cost=4.2;

ii. WeightedMultipathOpt

Input string: 1100 Output string: 0000 ;cost=2.2; Output string: 0000 ;cost=3.2;

Fewer paths because we are taking out redundant paths.

iii. *T_{out}*: gives all of T's outputs. *xT_{out}*: gives all of x@T's outputs. *Ty_n*: gives all of T@y's possible inputs. *xTy*: T accepts x as input and output y. *exTye*: turns the empty input string into the empty output string.

exTye_{opt}, since it's optimizing the FSM that turns ϵ to ϵ , should only have two states and one arc.

The last three FSMs are important because they are optimizing components of the final FSM. *xT_{out}_{opt}* optimizes the *x@T* part, and *Ty_n_{opt}* optimizes the *T@y*.

9 Q10

9.1 b

The sentences generated from ngramrandgen are anywhere from 2 words long to 30-ish words long.

fstshortestpath gives us only one output, ".". This is because it's not in the training data but is backed off. This is the shortest path, thus, it's given as output.

9.2 d

"Andy cherished the barrels each house made ."

Output string: Andycherishedthebarrelseachhousemade. jcost=67.4284j
Output string: Andycherishedthebarrelseachhousemade. jcost=69.3358j
Output string: Andycherishedthebarrelseachhousemade. jcost=79.479j
Output string: Andycherishedthebarrelseachhousemade. jcost=81.3864j

"If only the reporters had been nice ." Output string: Ifonlythereporter-shadbeennice. jcost=49.2604j
Output string: Ifonlythereportershadbeennice. jcost=49.5864j
Output string: Ifonlythereportershadbeennice. jcost=50.0427j
Output string: Ifonlythereportershadbeennice. jcost=50.3688j
Output string: Ifonlythereportershadbeennice. jcost=55.2341j ... There were a total of 31 output strings for this second input string.

"Thank you" Unable to parse input string.

Each input string with each of its potential output strings are pairs that are ends to paths in the FST that would accept these pairs. The differing costs are a result of using backoff (or not using it).

10 Q11

See noisy.grm please!

10.1 d

The results of the first sentence is correct and as expected. The second sentence failed.

The third sentence's parse shows that there are several ways to divvy up the spaces between certain words. Between "reporter said", we can have that or "reporters aid". With "is killed", we can have that or "i skilled". Furthermore, "everyone" can turn into "every one". Thus, we see sentences with every combination of these words.

The fourth sentence also failed.

The lowest cost sentence is included in this README. The highest cost the first sentence got to was 75.0509, and the highest cost the third sentence got to was 82.9002.

Input string: Ifonlythereporterhadbeennice. Output string: If only the reporter had been nice . jcost=50.2658j

Input string: If only. Rewrite failed.

Input string: ThereportersaidtothecitythatEveryoneIskilled. Output string:
 The reporters aid to the city that Everyone Is killed . jcost=70.201j
 Input string: Thankyou. Rewrite failed.

10.2 e

i. $w1$ = cost of each extra random character. $w2$ = minimum cost for each word generated.

iii. Total sum of probabilities = 1.

iv. Would lower $w1$, so we can have more extra characters for cheaper.

v. The probability of the random words would increase because not only do we get more extra random characters, we also see the minimum cost of each word decrease. This means that the decoder would likely rip words apart (for example, turning "everyone" into "every" and "one"). This may mean the decoder doesn't work as well as it should.

vi. We can use an NGram model on the FSA.

10.3 f

The results are varied. Sometimes we can get expected results, where the junks symbol will replace full words that are OOV. In other cases, however, junks can rip words apart. "Everyone" can turn into "junks" but it can also turn into "junks one", or "junks very junks", which isn't quite what we want.

10.4 g

junks is still ripping words apart. "recommending" turned into "junks ending".

We could increase the weight of $w2$ in order to ensure that words are more likely to be in the vocabulary.

10.5 h

Please run:

editdist entest.txt entest-recovered.txt

Output: Total edit distance 709 over 50 lines (about 14 per line).

11 Q12

Used overleaf.com to generate LaTeX document.